

Lecture 18:

Interconnection Networks

CMU 15-418: Parallel Computer Architecture and Programming (Spring 2012)

Announcements

- **Project deadlines:**
 - **Mon, April 2: project proposal: 1-2 page writeup**
 - **Fri, April 20: project checkpoint: 1-2 page writeup**
 - **Thurs, May 10: final presentations + final writeup**

Today's Agenda

- **Interconnection Networks**
 - **Introduction and Terminology**
 - Topology
 - Buffering and Flow control

Interconnection Network Basics

- **Topology**

- Specifies way switches are wired
- Affects routing, reliability, throughput, latency, building ease

- **Routing**

- How does a message get from source to destination
- Static or adaptive

- **Buffering and Flow Control**


- What do we store within the network?
 - Entire packets, parts of packets, etc?
- How do we manage and negotiate buffer space?
 - How do we throttle during oversubscription?
- Tightly coupled with routing strategy

Terminology

- **Network interface**
 - Connects endpoints (e.g. cores) to network.
 - Decouples computation/communication
- **Links**
 - Bundle of wires that carries a signal
- **Switch/router**
 - Connects fixed number of input channels to fixed number of output channels
- **Channel**
 - A single logical connection between routers/switches

More Terminology


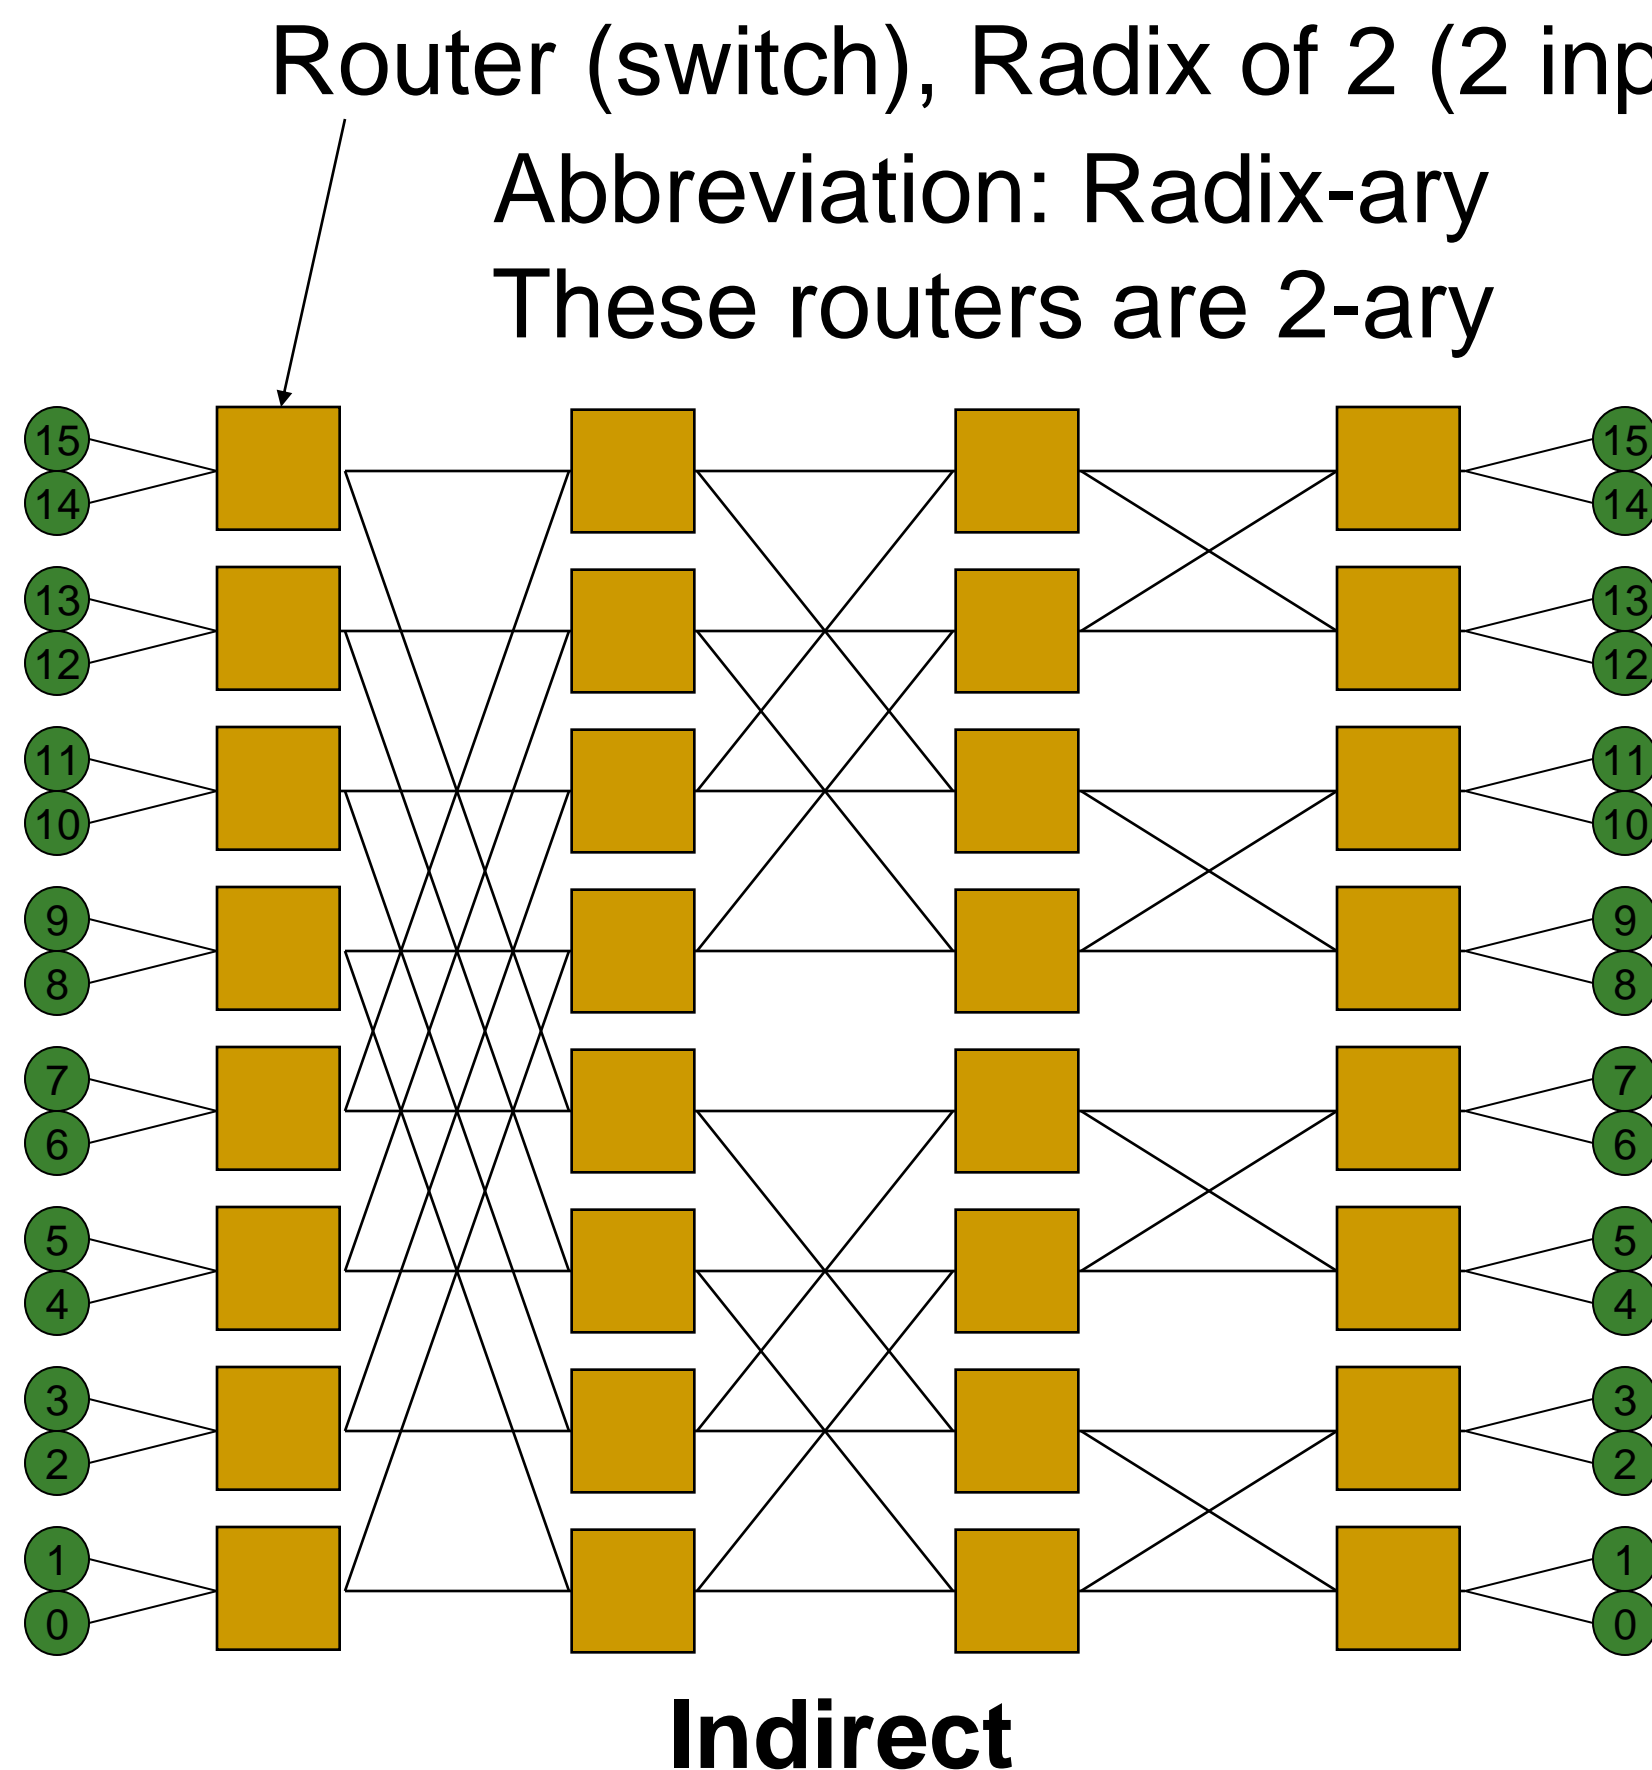
- **Node**
 - A network endpoint connected to a router/switch
- **Message**
 - Unit of transfer for network clients (e.g. cores, memory)
- **Packet**
 - Unit of transfer for network
- **Flit**
 - Flow control digit
 - Unit of flow control within network



Some More Terminology

■ Direct or Indirect Networks

- Endpoints sit “inside” (direct) or “outside” (indirect) the network
- E.g. mesh is direct; every node is both endpoint and switch



Today's Agenda

- **Interconnection Networks**
 - Introduction and Terminology
 - **Topology**
 - Buffering and Flow control

Properties of a Topology/Network

- **Regular or Irregular**
 - regular if topology is regular graph (e.g. ring, mesh)
- **Routing Distance**
 - number of links/hops along route
- **Diameter**
 - maximum routing distance
- **Average Distance**
 - average number of hops across all valid routes

Properties of a Topology/Network

■ Bisection Bandwidth

- Often used to describe network performance
- Cut network in half and sum bandwidth of links severed
 - (Min # channels spanning two halves) * (BW of each channel)
- Meaningful only for recursive topologies
- Can be misleading, because does not account for switch and routing efficiency

■ Blocking vs. Non-Blocking

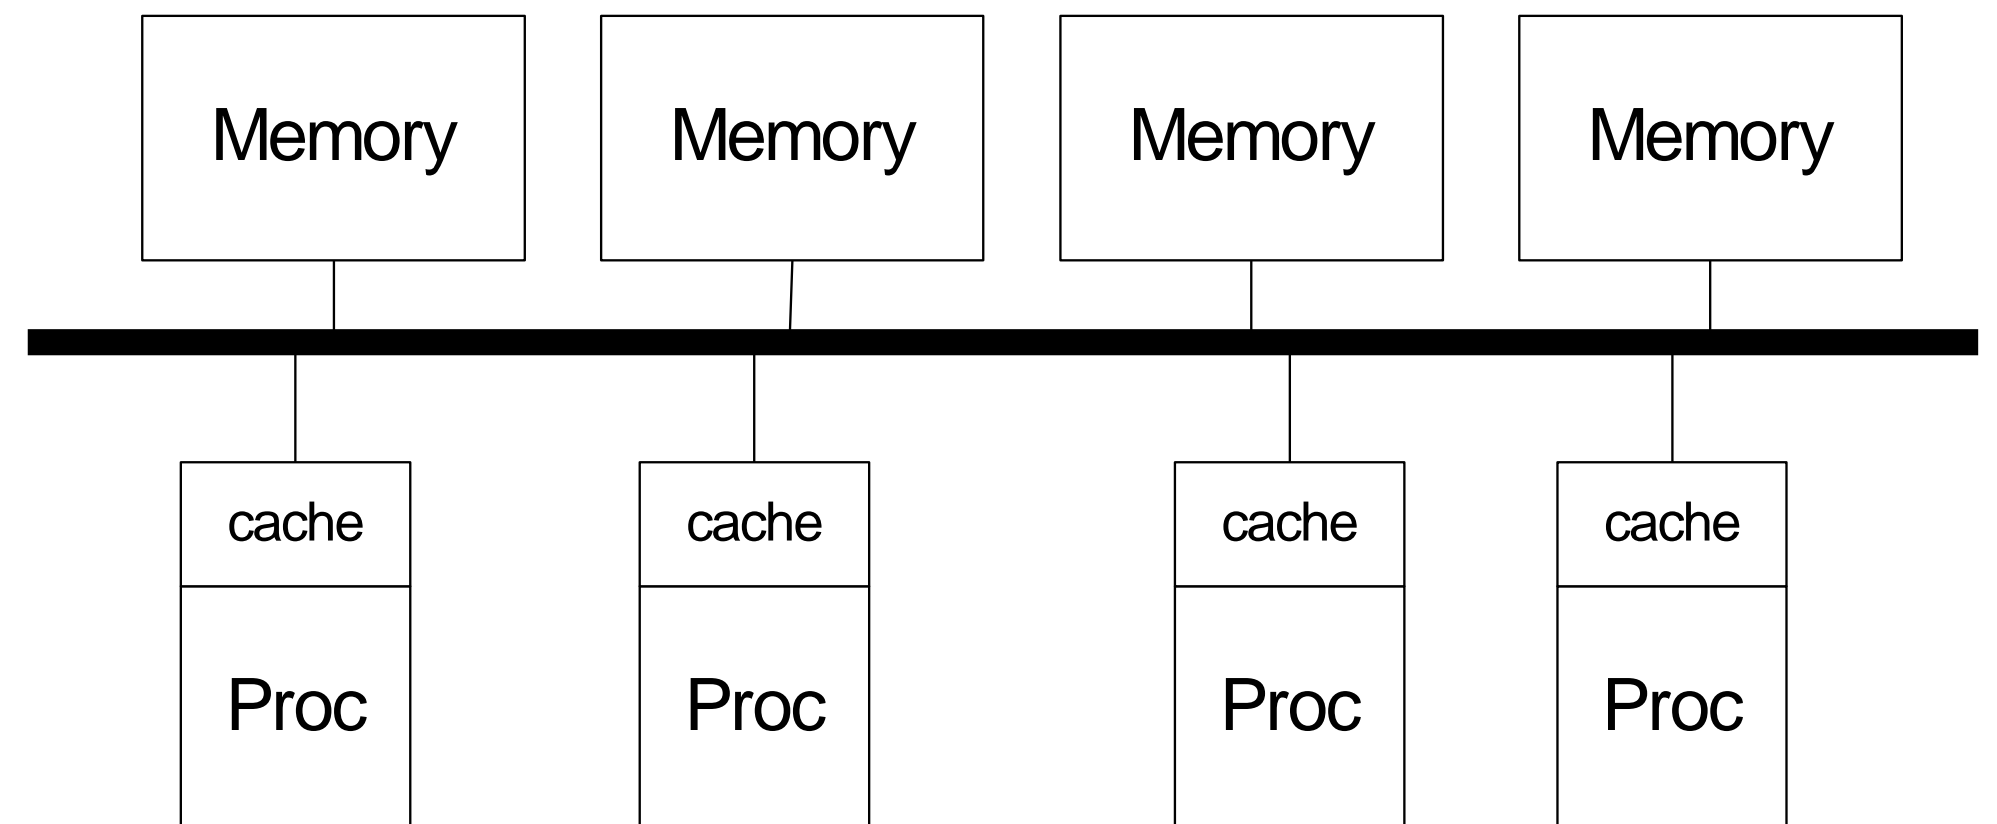
- If connecting any permutation of sources & destinations is possible, network is non-blocking; otherwise network is blocking.

Many Topology Examples

- **Bus**
- **Crossbar**
- **Ring**
- **Tree**
- **Omega**
- **Hypercube**
- **Mesh**
- **Torus**
- **Butterfly**
- **...**

Bus

- + Simple
- + Cost effective for a small number of nodes
- + Easy to implement coherence (snooping)
- Not scalable to large number of nodes
(limited bandwidth, electrical loading → reduced frequency)
- High contention




Crossbar

- **Every node connected to all others** (non-blocking)
- **Good for small number of nodes**
- + **Low latency and high throughput**
- **Expensive**
- **Not scalable $\rightarrow O(N^2)$ cost**
- **Difficult to arbitrate**

Core-to-cache-bank networks:

- **IBM POWER5**
- **Sun Niagara I/II**




Ring

- + Cheap: $O(N)$ cost
- High latency: $O(N)$
- Not easy to scale
 - Bisection bandwidth remains constant


Used in:

- Intel Larrabee/Core i7
- IBM Cell




Mesh

- **$O(N)$ cost**
- **Average latency: $O(\sqrt{N})$**
- **Easy to layout on-chip: regular & equal-length links**
- **Path diversity: many ways to get from one node to another**
- **Used in:**
 - **Tilera 100-core CMP**
 - **On-chip network prototypes**



Torus

- Mesh is not symmetric on edges: performance very sensitive to placement of task on edge vs. middle
- Torus avoids this problem
- + Higher path diversity (& bisection bandwidth) than mesh
- Higher cost
- Harder to lay out on-chip
- Unequal link lengths



Trees

Planar, hierarchical topology

Latency: $O(\log N)$



Good for local traffic

+ Cheap: $O(N)$ cost

+ Easy to Layout

- Root can become a bottleneck

Fat trees avoid this problem (CM-5)



Hypercube



- Latency: $O(\log N)$
- Radix: $O(\log N)$
- #links: $O(N \log N)$

+ Low latency

- Hard to lay out in 2D/3D


▪ Used in some early message passing machines, e.g.:

- Intel iPSC
- nCube






Multistage Logarithmic Networks

- **Idea:** Indirect networks with multiple layers of switches between terminals
- **Cost:** $O(N \log N)$, **Latency:** $O(\log N)$
- **Many variations** (Omega, Butterfly, Benes, Banyan, ...)
- **E.g. Omega Network:**



Q: Blocking or non-blocking?

Review: Topologies



Topology	Crossbar	Multistage Logarith.	Mesh
Direct/Indirect	Indirect	Indirect	Direct
Blocking/ Non-blocking	Non-blocking	Blocking	Blocking
Cost	$O(N^2)$	$O(N \log N)$	$O(N)$
Latency	$O(1)$	$O(\log N)$	$O(\sqrt{N})$


Today's Agenda

- **Interconnection Networks**
 - Introduction and Terminology
 - Topology
 - **Buffering and Flow control**

Circuit vs. Packet Switching

- **Circuit switching sets up full path**

- Establish route then send data
- (no one else can use those links)
- faster and higher bandwidth
- setting up and bringing down links slow




- **Packet switching routes per packet**

- Route each packet individually (possibly via different paths)
- if link is free can use
- potentially slower (must dynamically switch)
- no setup, bring down time




Packet Switched Networks: Packet Format

- **Header**
 - routing and control information
- **Payload**
 - carries data (non HW specific information)
 - can be further divided (framing, protocol stacks...)
- **Error Code**
 - generally at tail of packet so it can be generated on the way out



Handling Contention



- **Two packets trying to use the same link at the same time**
- **What do you do?**
 - **Buffer one**
 - **Drop one**
 - **Misroute one (deflection)**
- **We will only consider buffering in this lecture**

Flow Control Methods


- **Circuit switching**
- **Store and forward (Packet based)**
- **Virtual Cut Through (Packet based)**
- **Wormhole (Flit based)**

Circuit Switching Revisited

- Resource allocation granularity is high
 - Idea: Pre-allocate resources across multiple switches for a given “flow”
 - Need to send a probe to set up the path for preallocation
-
- + No need for buffering
 - + No contention (flow’s performance is isolated)
 - + Can handle arbitrary message sizes
 - Lower link utilization: two flows cannot use the same link
 - Handshake overhead to set up a “circuit”

Store and Forward Flow Control


- Packet based flow control
- Store and Forward
 - Packet copied entirely into network router before moving to the next node
 - Flow control unit is the entire packet
- Leads to high per-packet latency
- Requires buffering for entire packet in each node



Can we do better?

Cut through Flow Control

- Another form of packet based flow control
- Start forwarding as soon as header is received and resources (buffer, channel, etc) allocated
 - Dramatic reduction in latency
- Still allocate buffers and channel bandwidth for full packets




- What if packets are large?

Cut through Flow Control

- What to do if output port is blocked?
- Lets the tail continue when the head is blocked, absorbing the whole message into a single switch.
 - Requires a buffer large enough to hold the largest packet.
- Degenerates to store-and-forward with high contention
- Can we do better?

Wormhole Flow Control

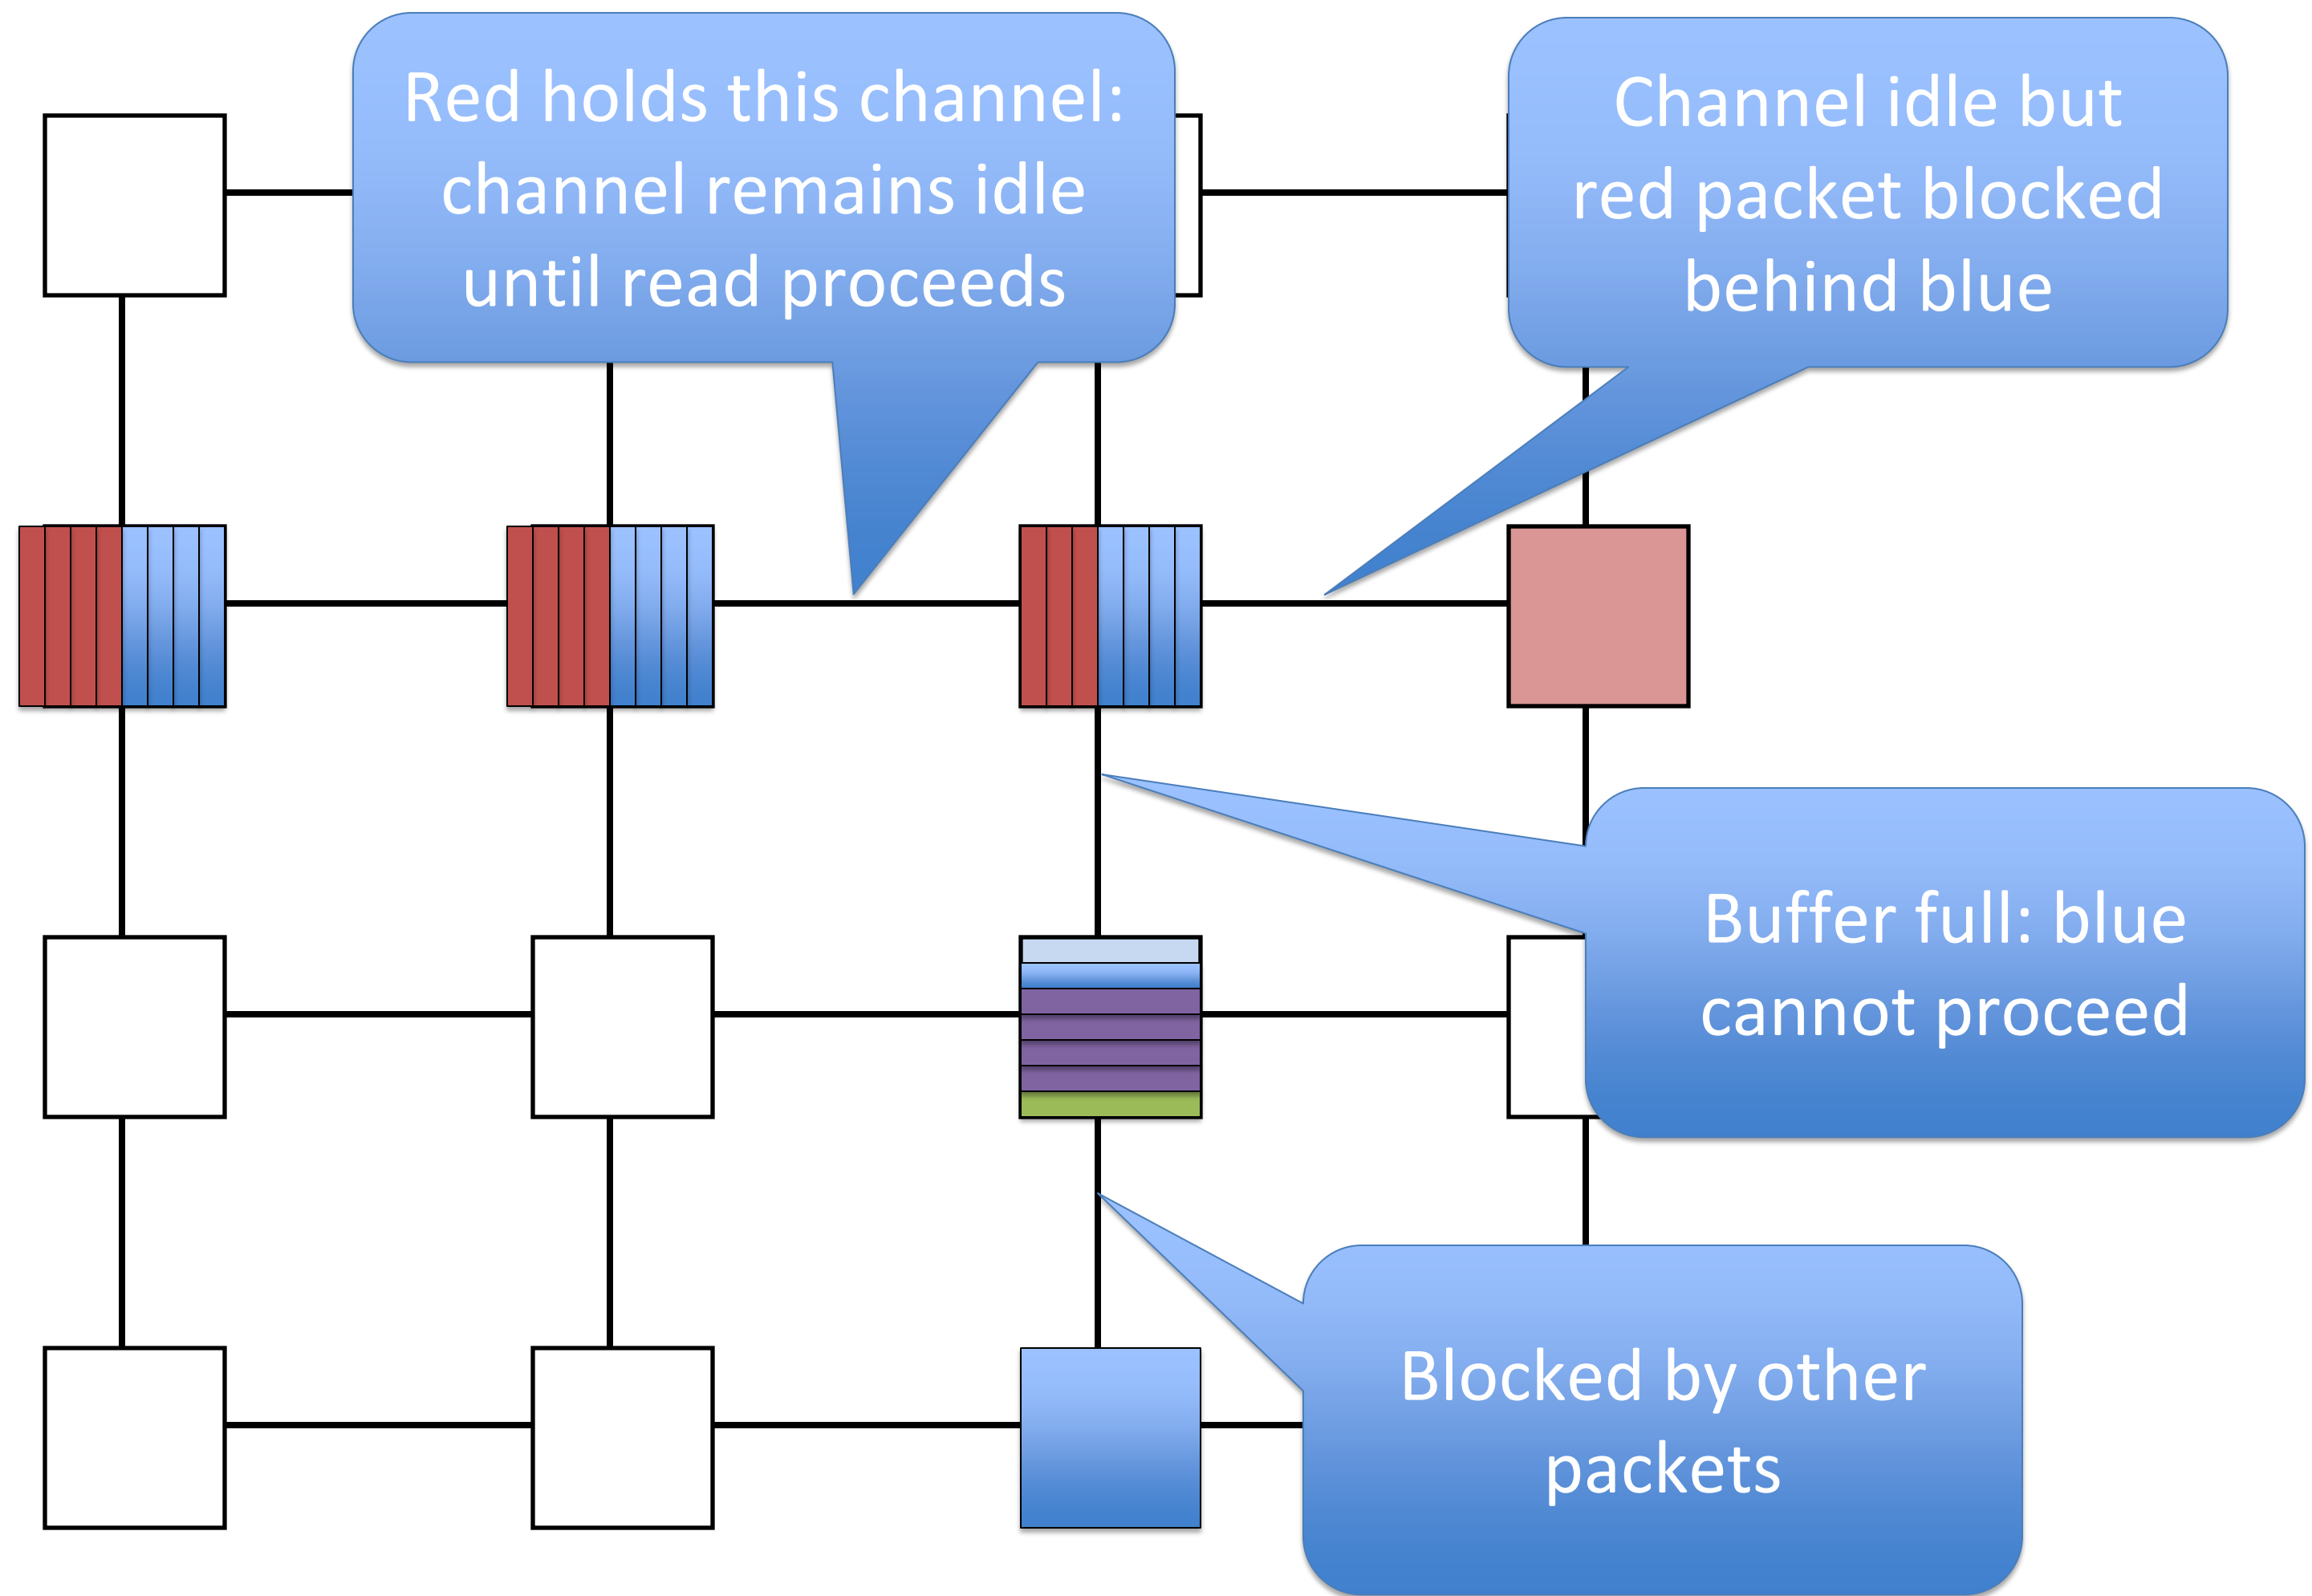


- Packets broken into (potentially) smaller flits (buffer/bw allocation unit)
- Flits are sent across the fabric in a *wormhole fashion*
 - Body follows head, tail follows body
 - Pipelined
 - If head blocked, rest of packet stops
 - Routing (src/dest) information only in head
- How does body/tail know where to go?
- Latency almost independent of distance for long messages

Wormhole Flow Control

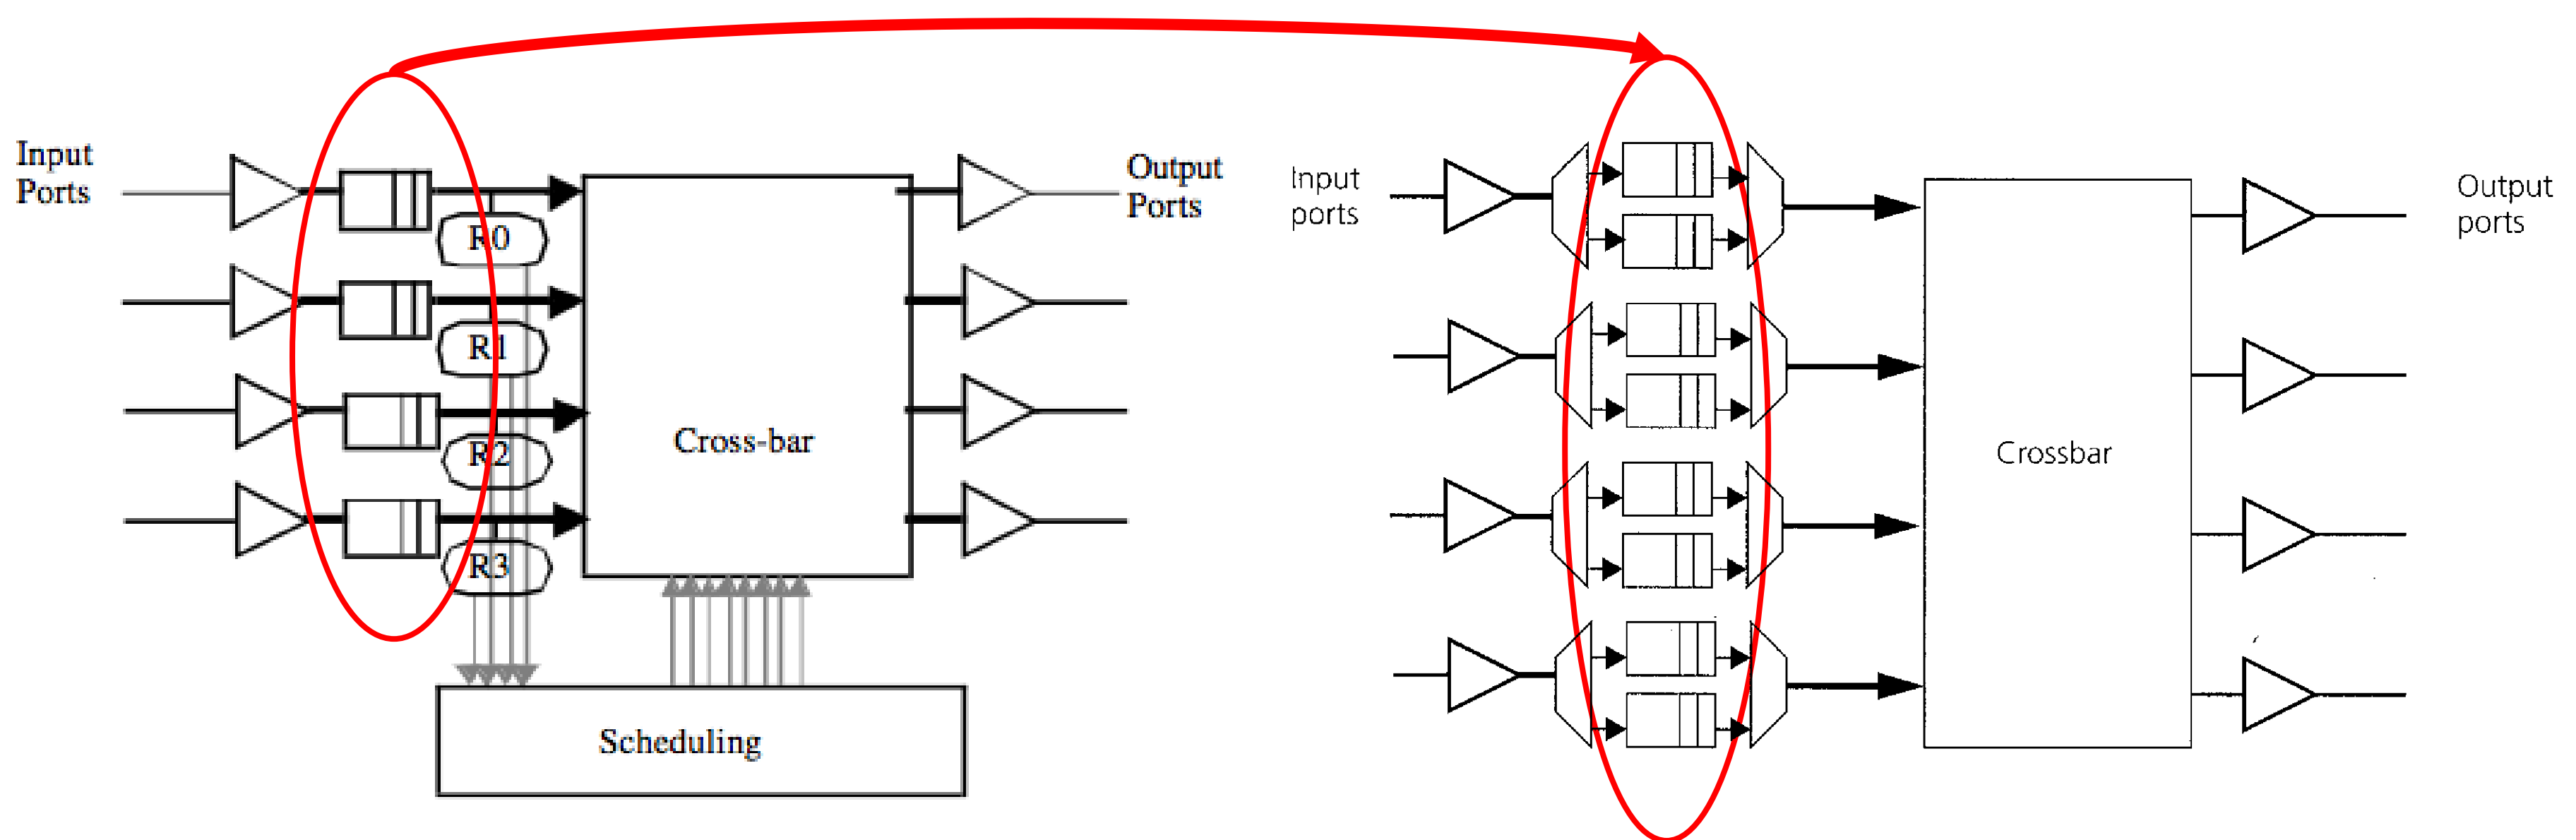
- Advantages over “store and forward” flow control
 - + Lower latency
 - + More efficient buffer utilization
- Limitations
 - Suffers from **head-of-line (HOL) blocking**
 - If head flit cannot move due to contention, another worm cannot proceed even though links may be idle

Head-of-Line Blocking

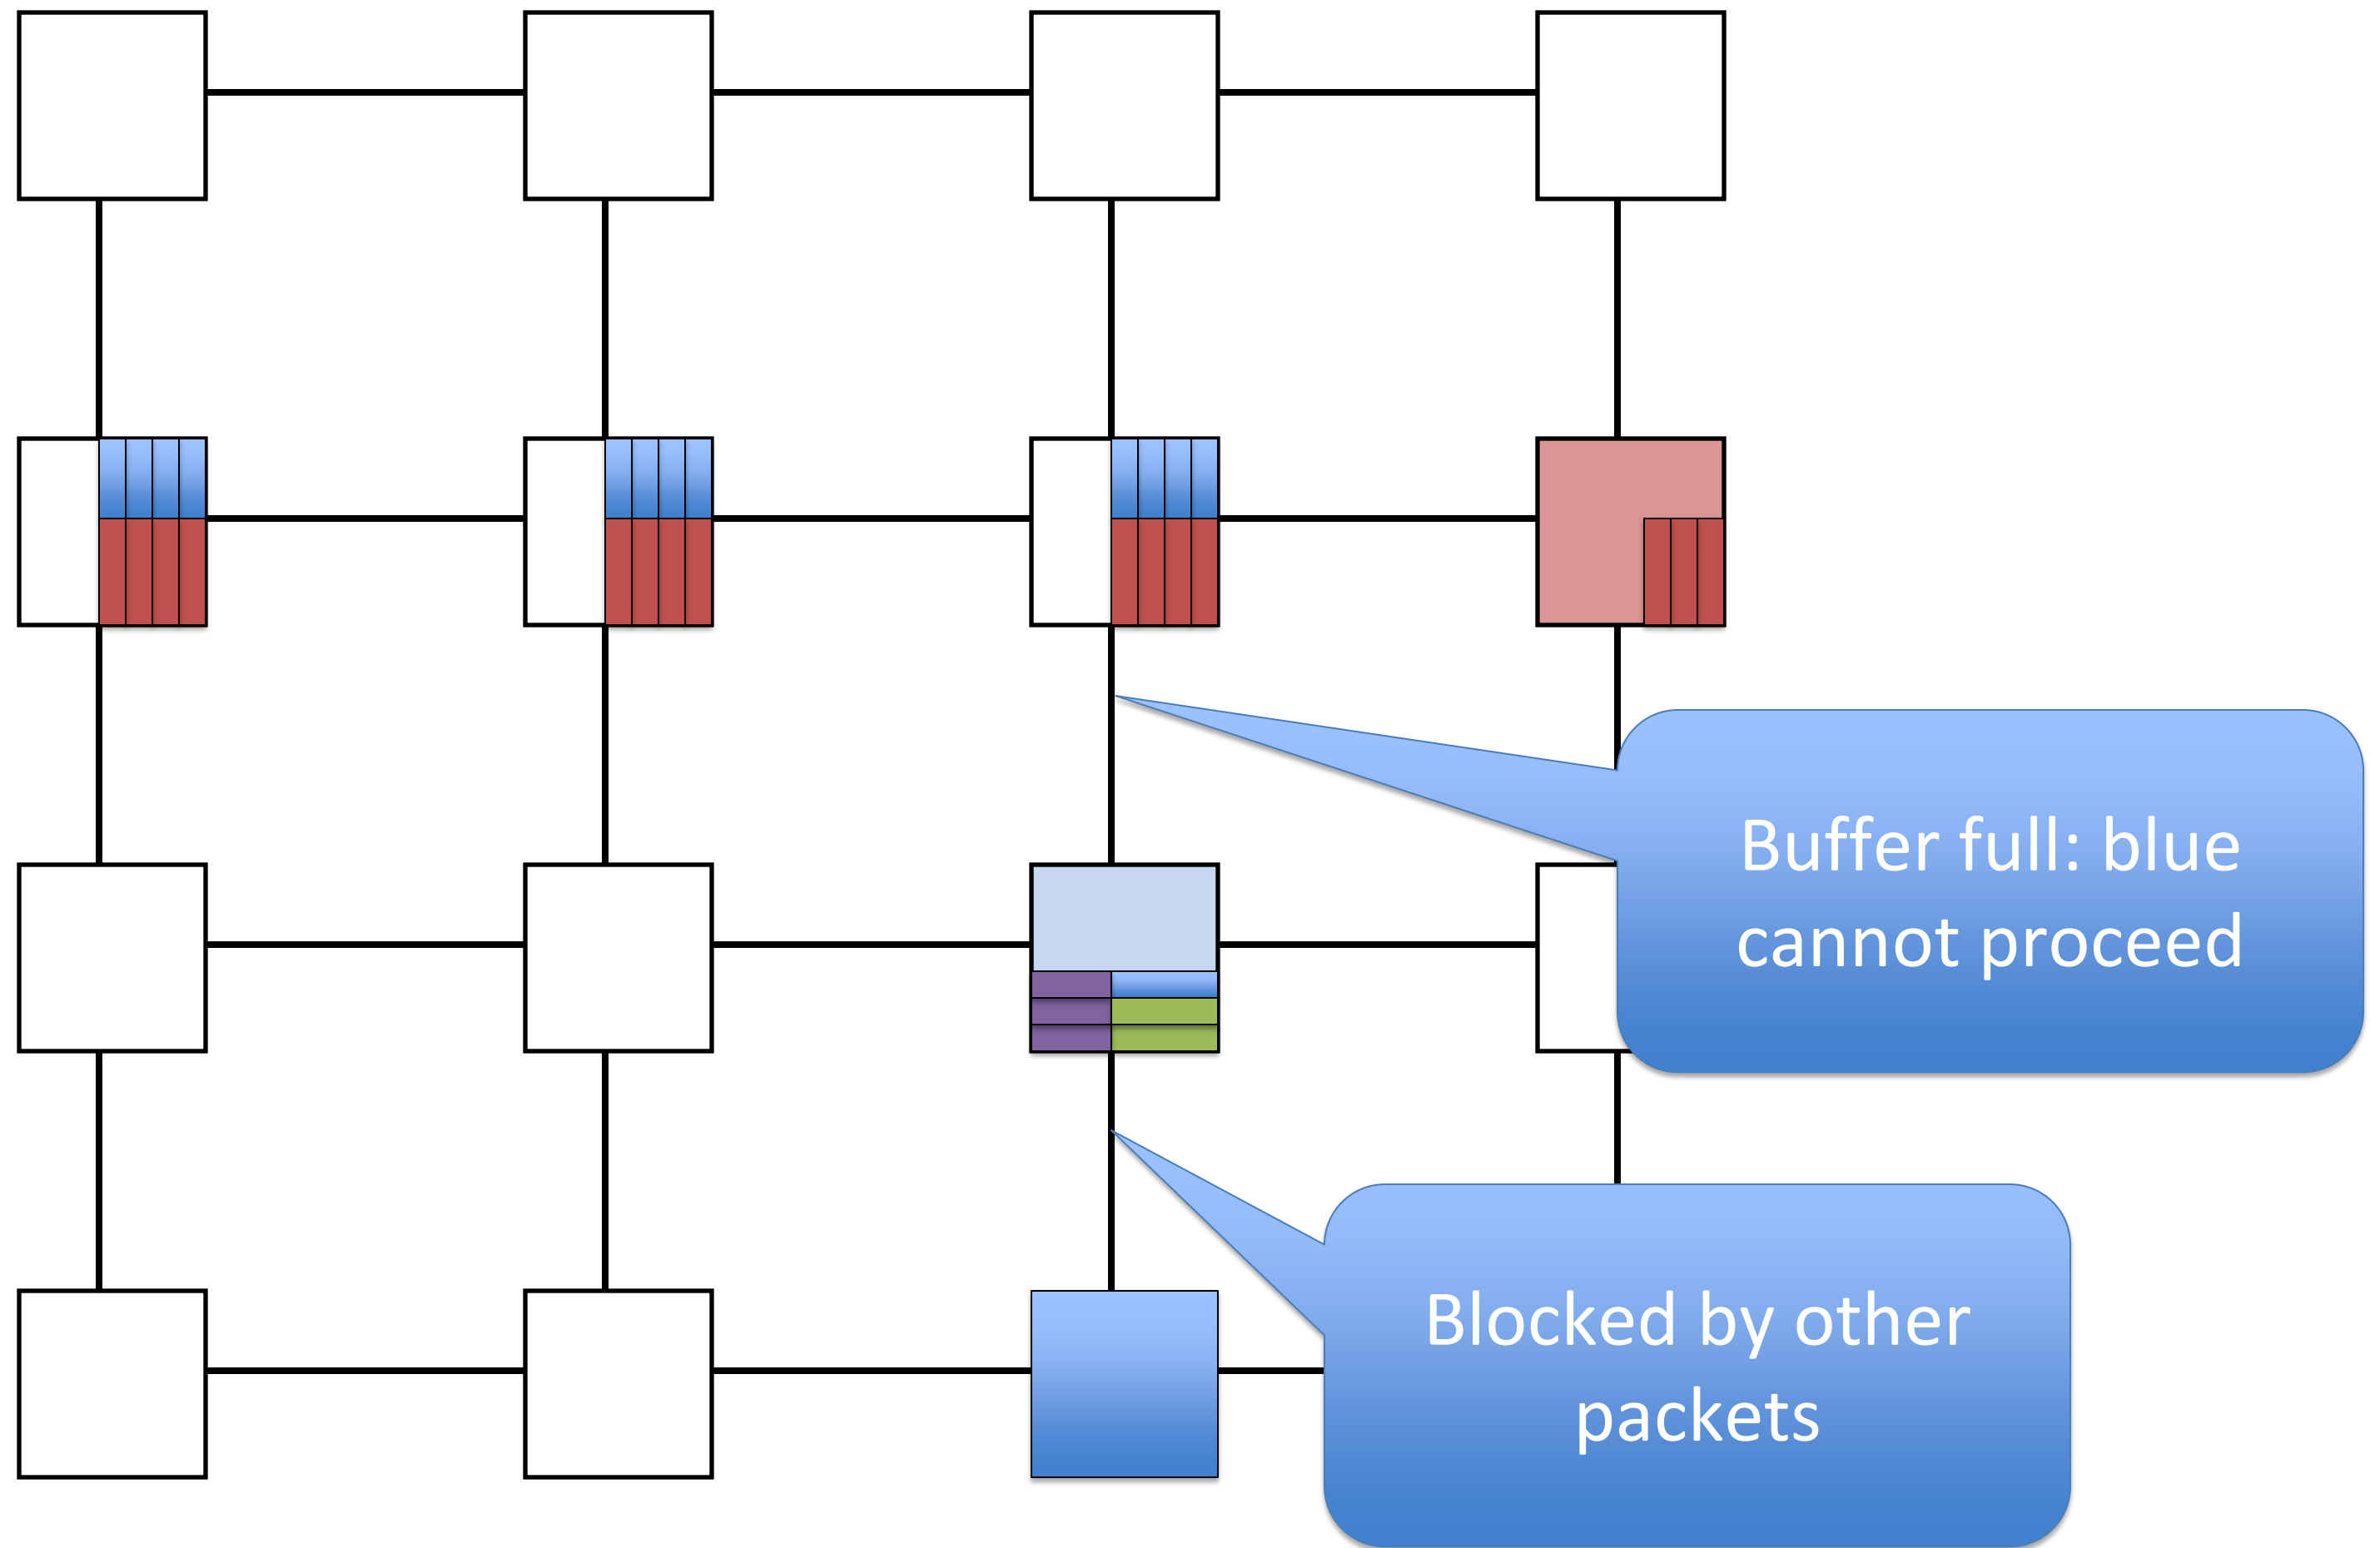


Virtual Channel Flow Control

- **Idea: Multiplex multiple channels over one physical channel**
- **Reduces head-of-line blocking**
- **Divide up the input buffer into multiple buffers sharing a single physical channel**
- **Dally, “Virtual Channel Flow Control,” ISCA 1990.**




Virtual Channel Flow Control



Other Uses of Virtual Channels

■ Deadlock avoidance

- Enforcing switching to a different set of virtual channels on some “turns” can break the cyclic dependency of resources
- Escape VCs: Have at least one VC that uses deadlock-free routing. Ensure each flit has fair access to that VC.
- Protocol level deadlock: Ensure request and response packets use different VCs → prevent cycles due to intermixing of different packet classes



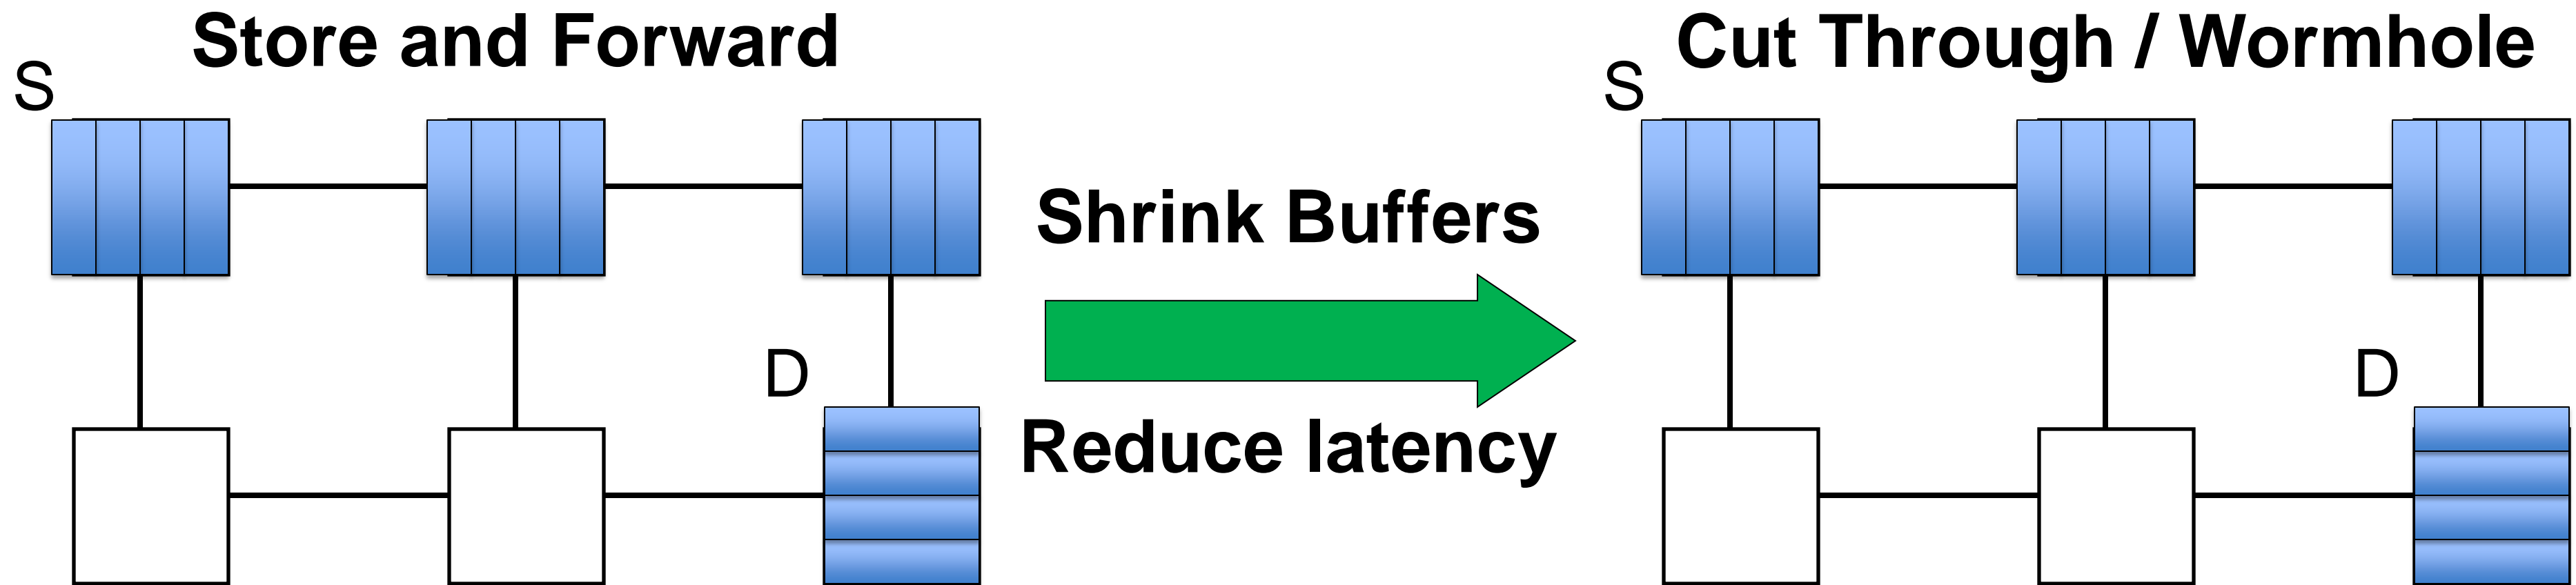
■ Prioritization of traffic classes

- Some virtual channels can have higher priority than others

Communicating Buffer Availability

- **Credit-based flow control**
 - Upstream knows how many buffers are downstream
 - Downstream passes back credits to upstream
 - Significant upstream signaling (esp. for small flits)
- **On/Off (XON/XOFF) flow control**
 - Downstream has on/off signal to upstream
- **Ack/Nack flow control**
 - Upstream optimistically sends downstream
 - Buffer cannot be deallocated until ACK/NACK received
 - Inefficiently utilizes buffer space

Review: Flow Control

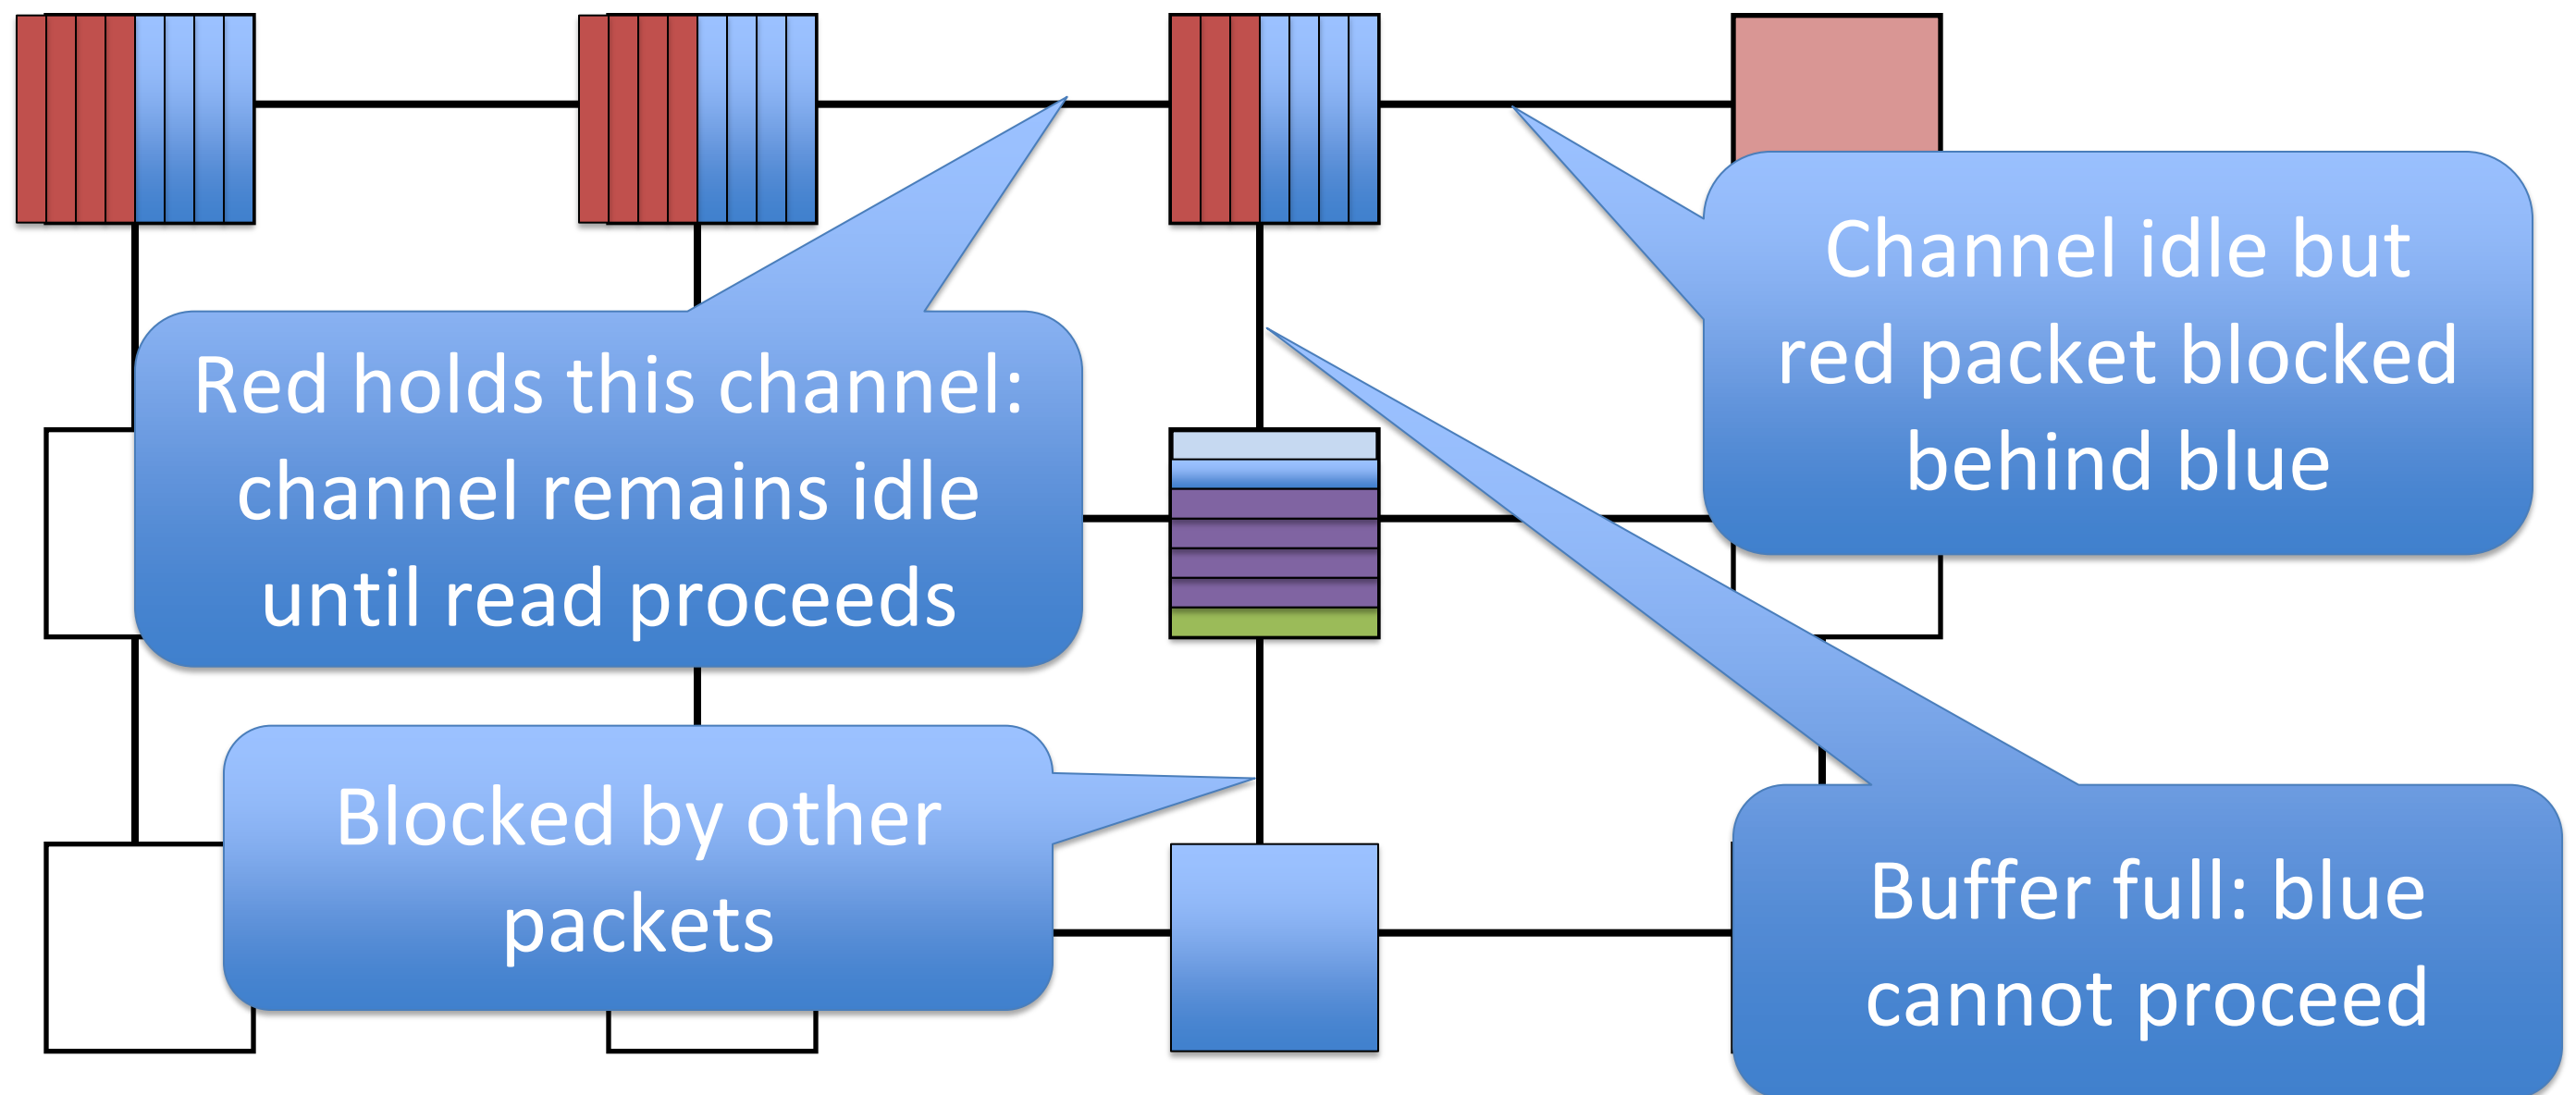


Any other
issues?

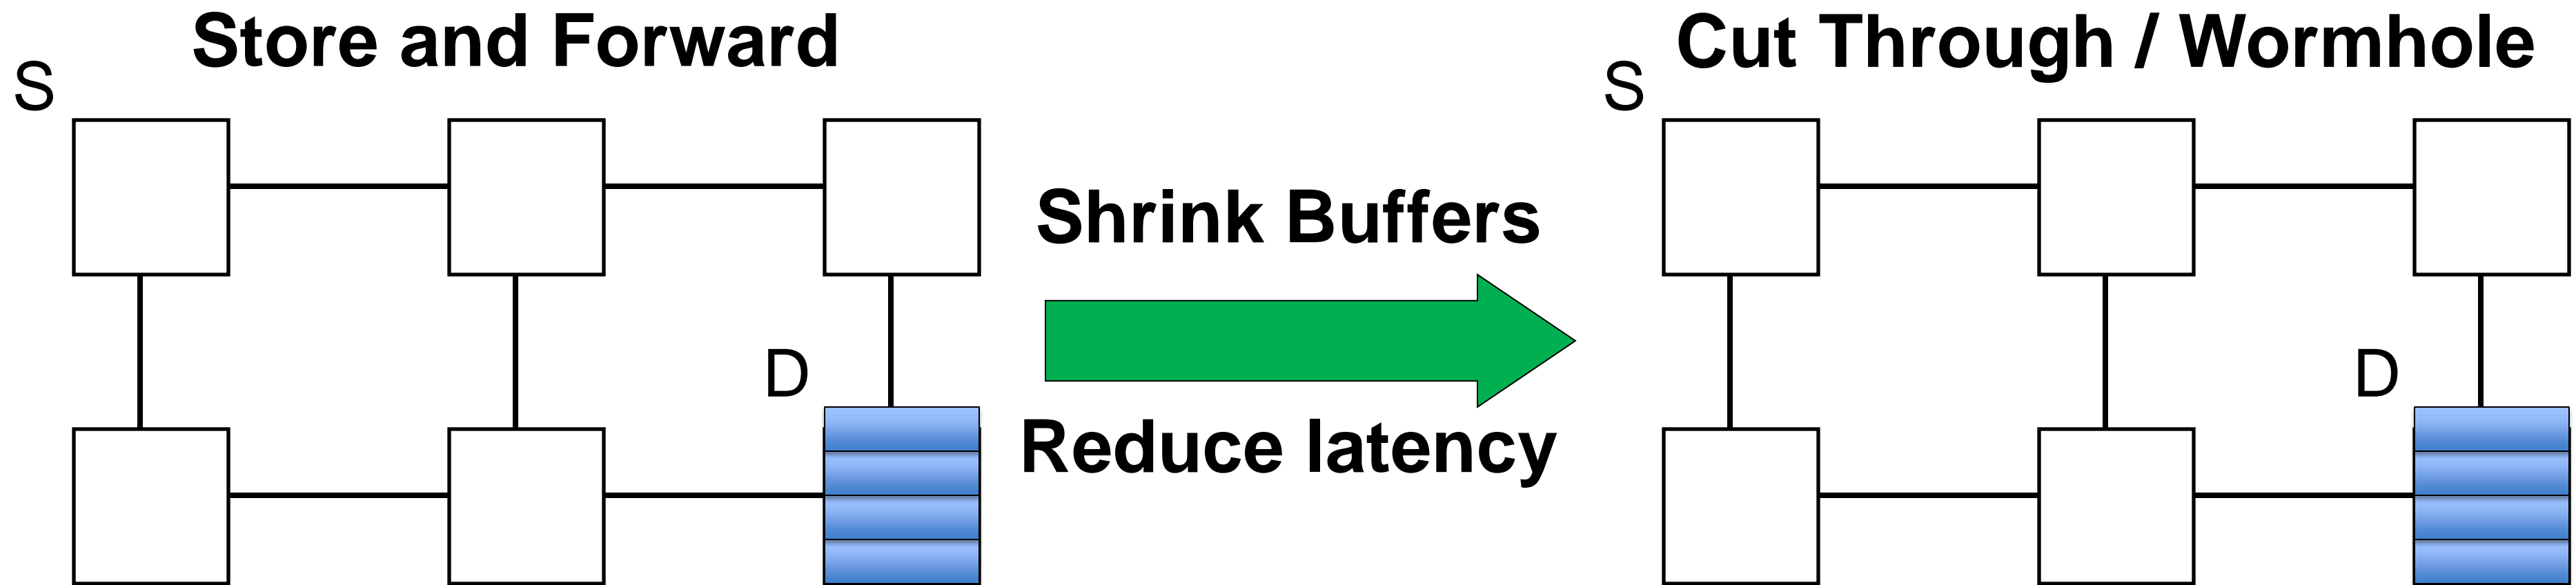
**Head-of-Line
Blocking**

↓

**Use Virtual
Channels**



Review: Flow Control



Any other
issues?

**Head-of-Line
Blocking**

↓

**Use Virtual
Channels**

